

Script that processes HL7 data from JSON files stored in S3 buckets or RedisJSON and stores the transformed/processed data in PostgreSQL schema.

[See LinkedIn Post](#)

Requires:

Use `create_sample_data.py` to generate anonymized sample data

1. Apache Spark Standalone or Cluster (3.4.1 or later)
I set up a Standalone Spark Master with 5 worker nodes
<https://spark.apache.org/downloads.html>
2. Have following JARS downloaded - in my case for Spark v3.4.1:
 `/var/tmp/sparkjars/postgresql-42.6.0.jar`
 `/var/tmp/sparkjars/aws-java-sdk-bundle-1.12.262.jar`
 `/var/tmp/sparkjars/hadoop-aws-3.3.4.jar`
3. Postgres JDBC driver (postgresql-42.6.0.jar or later)
4. S3 Bucket name (<s3_bucket_name>) (Mirth connect is storing
 RAW HL7 messages as JSON file
 in a S3 bucket)
5. S3 Bucket Prefix (yyyy/m/d)
6. S3 Bucket contains, single, or a set of JSON files that are
 formatted like the Sample JSON file.
7. etl.config file to store configuration details etc:
 [reportdb]
 host=
 port=5432
 dbname=
 dbuser=
 dbuserpass=

 [spark]
 master=
 masterport=7077

 [redis]
 host=
 port=6379

 [aws]
 access.key=
 secret.key=
 [constants]
 IGNORE_SEG_FIELDS=['PID_1','PID_12','PV1_1','IN1_1','EVN_1','OBX_1','AL1_
 IGNORE_COMPONENT_FIELDS=['CX_4','CX_5','XTN_2','XTN_3','XTN_5','XTN_6','X
 HL7_SEGMENTS=['pid','pv1','pv2','pd1','evn','in1','in2','obx','all','gt1'
8. List of fields to ignore stored in "hl7_field_names_to_ignore.txt" file
 Sample:

al1_1_set_id_al1_si_none_1
 evn_1_event_type_code_id_none
 evn_5_operator_id_xcn_2_family_name_2
 evn_5_operator_id_xcn_3_given_name_2
 evn_5_operator_id_xcn_9_assigning_authority_1
 gt1_10_guarantor_type_is_none_1
 gt1_11_guarantor_relationship_ce_1_identifier_1
 gt1_12_guarantor_ssn_st_none_1
 gt1_16_guarantor_employer_name_xpn_1_family_name_1
 gt1_17_guarantor_employer_address_xad_1_street_address_1
 gt1_17_guarantor_employer_address_xad_3_city_1
 gt1_17_guarantor_employer_address_xad_4_state_or_province_1
 gt1_17_guarantor_employer_address_xad_5_zip_or_postal_code_1
 gt1_17_guarantor_employer_address_xad_6_country_1
 gt1_18_guarantor_employer_phone_number_xtn_1_telephone_number_1
 gt1_20_guarantor_employment_status_is_none_1
 gt1_23_guarantor_credit_rating_code_ce_1_identifier_1
 gt1_2_guarantor_number_cx_1_id_number_1
 gt1_3_guarantor_name_xpn_4_suffix_e_g_jr_or_iii_1
 gt1_3_guarantor_name_xpn_5_prefix_e_g_dr_1
 gt1_3_guarantor_name_xpn_7_name_type_code_1
 gt1_4_guarantor_spouse_name_xpn_1_family_name_1
 gt1_4_guarantor_spouse_name_xpn_2_given_name_1
 gt1_5_guarantor_address_xad_1_street_address_1

9. Field mapping file "field_map.txt",

see field_name = f'{ac_name}_{ac_long_name}_{fc_name}_{fc_long_name}':

Sample:

('pid_2_patient_id_cx_1_id_number','pt_id_4')
 ('pid_30_patient_death_indicator_id_none','pt_death_indicator')
 ('pid_3_patient_identifier_list_cx_1_id_number_2','pt_id_2')
 ('pid_3_patient_identifier_list_cx_1_id_number_3','pt_id_5')
 ('pid_4_alternate_patient_id_pid_cx_1_id_number','pt_id_3')
 ('pid_5_patient_name_xpn_1_family_name','pt_last_name')

How-to Run - S3:

```
>./s3_json_to_psql_etl.py --adt-feed-name "acme,roadrunner" --s3-bucket-prefi:

23/09/02 16:15:08 WARN NativeCodeLoader: Unable to load native-hadoop library
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogL

**** Starting for acme

23/09/02 16:15:10 WARN MetricsConfig: Cannot locate configuration: tried hadoo
23/09/02 16:15:53 WARN package: Truncated the string representation of a plan

**** Stored data in table v4_acme
**** Completed for acme
**** Starting for roadrunner
**** Stored data in table v4_roadrunner
**** Completed for roadrunner
```

How-to Run - RedisJSON:

```
>./s3_redis_json_to_psql_etl.py --adt-feed-name redis_$(date '+%s')

23/10/05 10:22:27 WARN NativeCodeLoader: Unable to load native-hadoop library
Setting default log level to "WARN".

**** Starting for redis_1696526545
**** Get Redis Jsons ****
**** Lower Case Column Names ****
**** Process HL7 ****
**** Start Process ****
```

Example RedisJSON:

Sample RedisJSON Key Names:

Total: 1 010 555

JSON

pid_8537764_d6c113a1058146faabc31b395ba3122b.json

JSON

pid_7260099_ffa85a85808249638e1cb5a9bab12a15.json

JSON

pid_5762946_d36c3432552b4dd7bb6f76e4fc4d3fe9.json

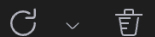
Sample RedisJSON Key Content:

JSON

pid_8537764_d6c113a1058146faabc31b395ba3122b.json

Key Size: 697 B Length: 609 TTL: No limit

Last refresh: < 1 min



```
{"patientId": "8537764", "tenantId": "tephromalacia", "dob": "1984-02-16", "id": "8537764", "updatedAt": "2023-01-10 17:44", "createdAt": "2023-01-10 17:44", "visitNumber": "52444378", "MSH": "MSH|^~&|EPICCARE|WB^WBPC|||20230110144357|S67031|ADT^A08^ADT_A01|400815517|P|2.3", "EVN": "EVN|A08|20230110144357||REGCHECKCOMP_A08|S67031^ROBINSON^CASEY^ANAME^^^^WB^^^^WBPC|WBPC^1740348929^SOMENAME", "PID": "PID|1||14891584^^^^EPI-62986117^^^^SOMERN||JOSHUA^FERNANDEZ||19840216|M|||123 MAIN ST^^SAN CITY^IL^12345^USA^P^^SC", "PV1": "PV1|||WBPCIVTA^^WBPC^^WBPC^^WBPC INFUSION THERAPY|||||||||||||1219753048"}"
```

Example table names:

```
public | v4_acme          | table |
public | v4_roadrunner         | table |
```

Sample JSON file:

```
{
  "patientId": "123456789",
  "tenantId": "<tenant id>",
  "dob": "1990-01-01",
  "id": 123456788,
  "updatedAt": "2023-09-01 00:59:22",
  "createdAt": "2023-09-01 00:59:22",
  "visitNumber": "987654321",
  "MSH": "MSH|^~&|EPIC CARE|WB^WBVC|||20230901095922|S327628|ADT^A08^ADT_A01|7",
  "EVN": "EVN|A08|20230901095922||REG_UPDATE_VISIT_CHANGE|S327628^RN|LASTNAME^",
  "PID": "PID|1||12345678^^^^EPI~123456789^^^^SOMEMRN||FIRSTNAME^LASTNAME||19",
  "PV1": "PV1|||WBVCERDA^ED07^07^WBVC^^WBVC^^WBVC EMERGENCY|||||||1"
}
```

Sample PostgreSQL table:

Column	Type	Collation	Nullable	Default	Source
dob	text				event
pt_event_recorded_date_time_1	text				event
pt_event_reason_code	text				event
pt_opr_assigning_facility_1	text				event
pt_event_operator_id_num_1	text				event
pt_event_operator_last_name	text				event
pt_event_operator_first_name	text				event
pt_opr_assigning_auth_1	text				event
pt_event_facil_namespace_id	text				event
pt_event_facil_universal_id	text				event
pt_event_facil_universal_id_type	text				event
patientid	text				event
pt_address_st_1	text				event
pt_address_city	text				event
pt_address_state_prov	text				event
pt_address_zip_postal	text				event
pt_address_country	text				event
pt_address_type	text				event
pt_county_parish	text				event
pt_last_name	text				event
pt_first_name	text				event
pt_dob	text				event
pt_gender	text				event
pt_visit_number	text				event
pt_loc_poc_1	text				event
pt_loc_room_1	text				event
pt_loc_bed_1	text				event
pt_loc_facility_1	text				event
pt_loc_description	text				event
pt_event_occur_date_time	text				event

Phase 1 Workflow

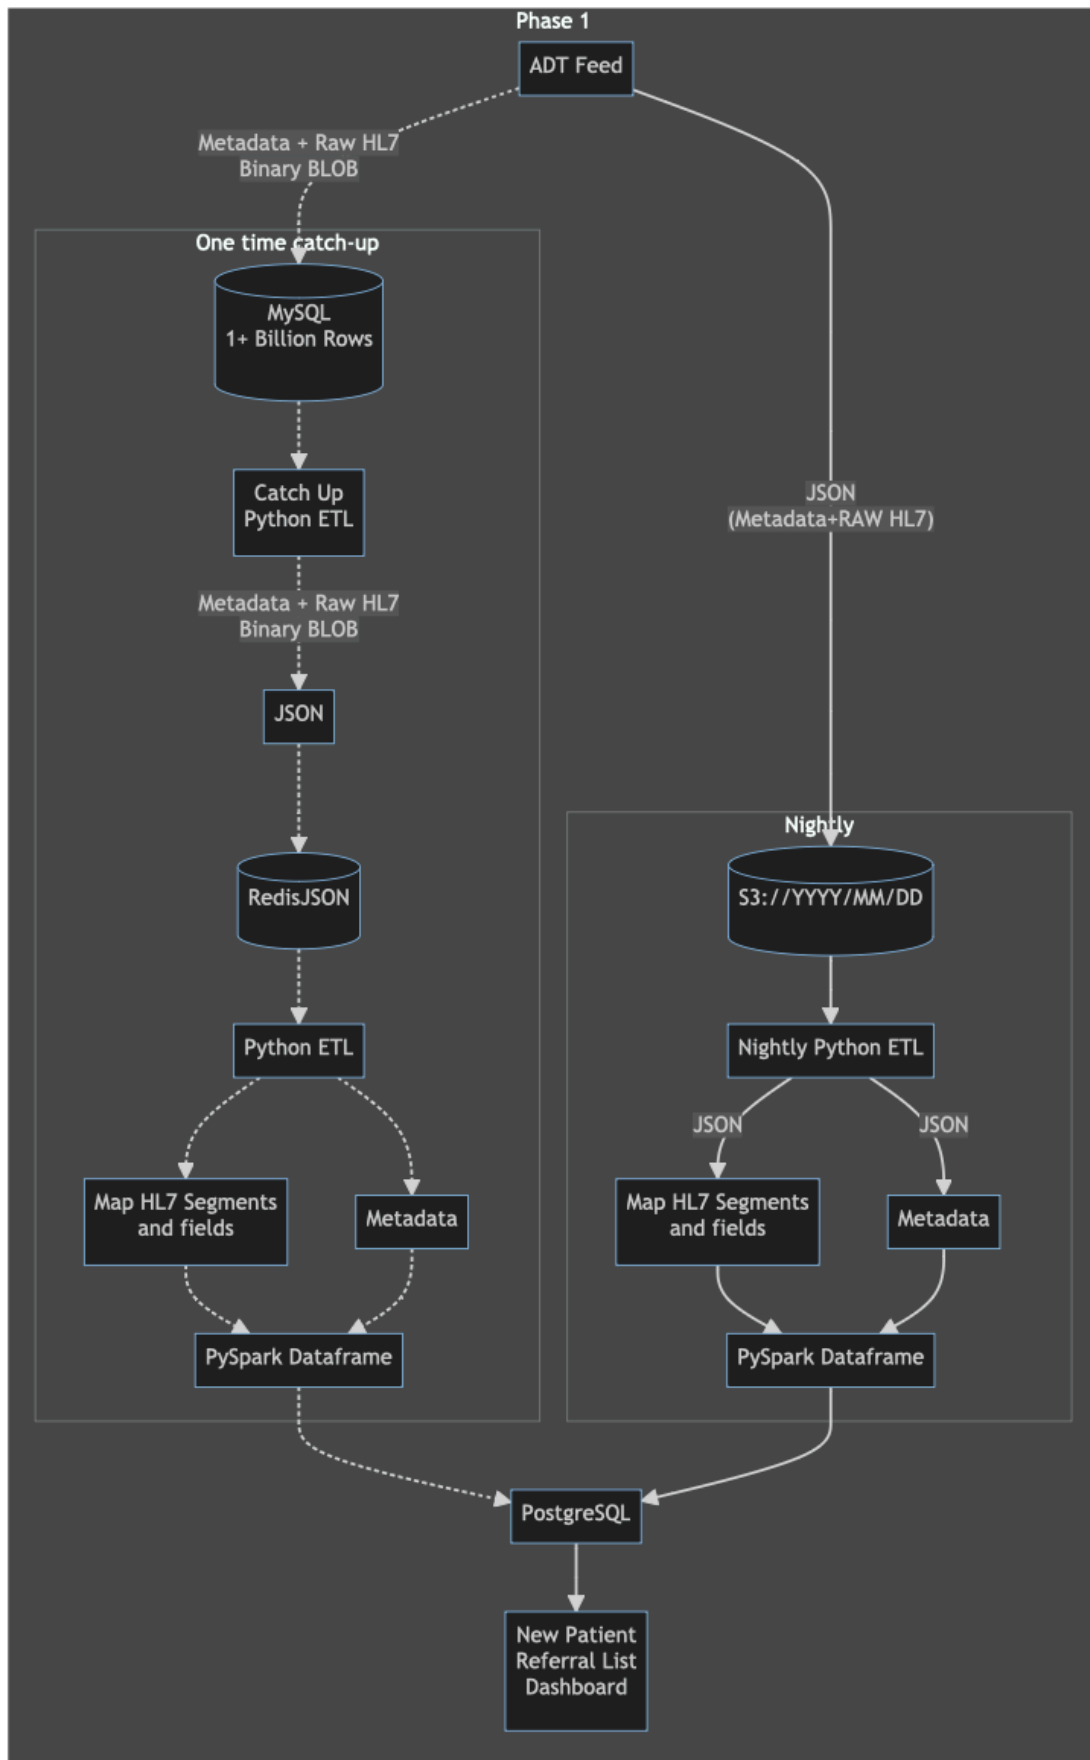
To catch up with over one billion rows (equivalent to three years' worth of raw data) in a MySQL Database and to facilitate the ingestion and processing of a large number of JSON files, I converted the rows into JSON files and stored each file in an S3 bucket as well as in the RedisJSON document store. The S3 bucket serves as permanent storage, while RedisJSON, due to

its

significantly faster read capabilities, was utilized for reading and processing tens of thousands of files simultaneously. Despite the initial cost of writing to both S3 and RedisJSON with over one billion rows, the reads, when comparing Python Boto3 with Python Redis, from RedisJSON

were exceptionally fast (.8 second for Boto3 , vs .007 seconds for PyRedis for a 3kb file).

The read size was only limited by the available RAM of the EC2 instance.

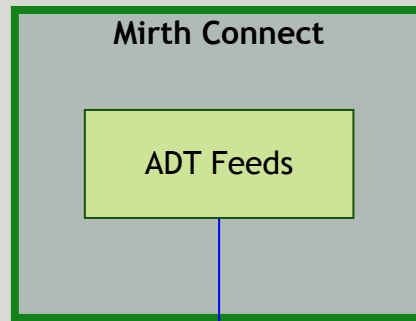


Final Implementation

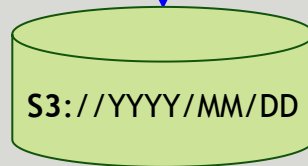
Since this has been my very first experience with processing HL7 data from ADT feeds, now that I have

learned quite a lot, my final implementation would look something like this.

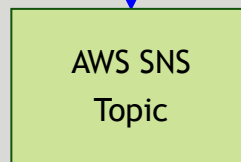
ADT ETL --> Data
Lakehouse



JSON
(Metadata
RAW HL7)



S3 Event
Notification



Publish
Message

